
THE GENETIC BASIS OF FITNESS IN *Drosophila*: A GENOME-WIDE ASSOCIATION STUDY

A PREPRINT

Heidi Wong

Department of Computer Science
Cranberry-Lemon University
Pittsburgh, PA 15213

Luke Holman *

School of Applied Sciences
Edinburgh Napier U
Edinburgh, UK.
l.holman@napier.ac.uk

April 6, 2021

Abstract

Enter the text of your abstract here.

Keywords blah · blee · bloo · these are optional and can be removed

Introduction

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

the *Drosophila* Genetic Reference Panel (DGRP), a collection of almost entirely homozygous lines that represent a snapshot of natural genetic variation from a population in North Carolina (REF: Mackay et al., 2012)

Methods

Fly stocks and husbandry

Our study focused on 125 lines randomly selected from the DGRP. All flies were reared in 25mm vials with Hoff food medium (REF?), lightly sprinkled with dried yeast, at a temperature of 25°C. We verified the genotype of each DGRP line using the restriction-based assay PCR described in (Mackay et al., 2012) for the eight most diagnostic markers, and verified the genotypes of the lines prior to data collection. In addition to the DGRP, we used two stocks carrying the visible markers, *brown*¹ (*bw*¹) and *PFRT(w^{hs})G13 P*Ubi*-GFP.nls2R1 P*Ubi*-GFP.nls2R2* as mates and competitors for the DGRP flies; these stocks are hereafter termed *bw* and *GFP* respectively (*GFP*: green fluorescent protein).

Measuring line means for male and female early- and later-life fitness

Overview: We aimed to produce a holistic measure of the fitness of males and females from each DGRP line, both for young flies (i.e. those that had just reached reproductive maturity) and older individuals. We regard a ‘fit female’ *Drosophila* as one that can survive in mixed-sex conditions, compete for food, and produce large

*Previous address: Melbourne.

numbers of eggs that hatch. We regard a fit male as one that that can similarly survive and feed, and which can outcompete other males in pre- and post-mating sexual selection. We therefore measured four phenotypic traits on each DGRP line using standardised assays, which for brevity we term early-life and late-life male and female fitness (while recognising that these are phenotypes that potentially correlate with fitness, and not fitness itself).

We first highlight some features of the design of our fitness assays that influence how one should interpret the four measured phenotypes. Firstly, both assays were designed to foster social interactions and competition between adults, such that our fitness measures potentially depend on phenotypes with a social or competitive component (e.g. female resistance to male courtship, or competition over limited yeast food). Secondly, the focal flies and their same-sex competitors were not replaced if they died, such that our fitness measures incorporate variation in mortality as well as progeny production. Thirdly, we quantified fitness by counting the number of offspring hatched, which avoids confounding variation in adult fitness components with variation in larva-to-adult survival. Finally, the focal flies in our fitness assays were reared from first-instar larvae at a standardised density of 100 per vial, minimising non-genetic differences between lines. We also reared the *bw* and *GFP* individuals in a standardised fashion by placing 15 1- to 4-day-old mated females into yeasted vials, then collecting virgin offspring 36h later.

The fitness assays were run across nine blocks, and DGRP line 352 was included in every block, providing a reference point to help statistically estimate block effects on fitness. In each block, we measured the four phenotypes in 8-17 lines, plus the reference line, 352.

Female fitness assay: 5 females from the focal DGRP line were placed in a food vial (termed the ‘interaction vial’) with 15 *GFP* males and 10 *bw* females (all flies were 2- to 3-day-old virgins). After allowing the flies to interact for 3 days, the 5 DGRP females were moved to an ‘egg collection’ vial filled with grape-agar medium to oviposit for 24h, before being returned to the original interaction vial with the *GFP* and *bw* flies. The eggs were allowed to develop into first instar larvae, where were then counted, giving the data used to estimate measure of early-life female fitness. To measure female late-life fitness, we waited 8 days (tipping once into a fresh vial), then replaced the old *GFP* males with 15 new 2- to 3 days-old virgin *GFP* males, and waited a further 3 days. The DGRP females were then moved to egg collection vials for 24h, and we again counted the total number of 1st instar larvae that emerged.

To estimate the line mean for each female early- and late-life fitness, we fit a Bayesian multivariate generalised linear mixed model (GLMM, implemented in the R package *brms*), with the early- and late-life progeny counts as response variables and line, block, and group ID as crossed random effects (the random effects were assumed to have correlated effects on the two response variables, with the correlations being estimated from the data). The response variables followed a Poisson distribution with log link. We then used the fitted model to derive the posterior predicted values for the line mean for each trait. These predicted values, which correct the fitness estimates for block effects and account for pseudoreplication, were used in all downstream analyses of the female fitness line means. We derived predictions on the scale of linear predictor and scaled them to have a mean of zero and unit variance. We also used the GLMM to find the posterior intraclass correlation coefficient for the line random effect, i.e. the proportion of variation in female fitness that is explained by DGRP line.

Male fitness assay: To measure male fitness, we placed 5 males from the focal DGRP line in an interaction vial with 10 *GFP* males and 15 *bw* females (all flies were 2- to 3-day-old virgins). After allowing the flies to interact and mate for 3 days, the females were moved to an egg collection vial for 24h. The females were then transferred back to their original interaction vial with the DGRP males and *GFP* males. After allowing the collected eggs to develop for 24h, and a random sample of up to 200 first instar larvae was collected and scored for *GFP* presence/absence, giving the data used to estimate male early-life fitness. The flies were then left for 8 days, and were tipped into a fresh interaction vial once during this time. Then, when the DGRP and *GFP* males were approximately 14 days old, the *bw* females were replaced with 15 new 2- to 3-day-old virgin *bw* females, and the flies were left to interact for 3 days. The females were then placed in a new egg collection vial to oviposit for 24h, and we scored the presence/absence of *GFP* in up to 200 larvae from these vials 24h later to yield the data for the late-life male fitness estimate.

We similarly used a multivariate GLMM to estimate the line mean male fitness. The procedure was similar to the model used for females, except that the response variable was the proportion of offspring sired (modelled using the binomial family and logit link), and we additionally corrected for the number of live competitor *GFP* males that were present at the time the females were removed for egg collection (by including the number of competitors as a covariate). Thus, we assume that the *GFP* males died randomly with respect to the genotype of the DGRP males, and adjusted the fitness of each line accordingly. We again derived the

posterior predictions of the line means, and estimated the proportion of variance in siring success that was explained by line.

Estimating quantitative genetic parameters

We estimated parameters such as heritabilities and additive genetic (co)variances from the parameter estimates of a Bayesian multivariate random effects model implemented in **brms** (REF). The response variable was a matrix of the four phenotypes for each line (estimated in the previously mentioned GLMMs), and the predictor variable was line. The line-level random effects were assumed to covary between phenotypes, according to the covariances in the genomic relatedness matrix, which we estimated from their multilocus genotypes using **sommer** (REF). We took the between-line (co)variance estimates as measures of the genetic (co)variances, and calculated heritability as the proportion of variance explained by line. We also calculated the posterior differences between pairs of genetic correlations, to test hypotheses about how these correlations differ by sex/age.

Univariate GWAS and multivariate adaptive shrinkage

In short, we performed a genome-wide association study (GWAS) using the software GEMMA (ref), and used multivariate adaptive shrinkage implemented in the R package **mashr** (REF) to process and interpret the GWAS results.

We included in the GWAS all loci (SNPs and indels) that had been genotyped in at least 90% of the DGRP lines in our study and which had a minor allele frequency of $\geq 5\%$. There were XX such loci, although only XX of these were in $<100\%$ linkage disequilibrium with all other loci; we therefore analysed these XX variants or variant groups (the latter were identified using PLINK). Missing genotypes were imputed using Beagle (REF). Using GEMMA’s **lmm** function, we ran one one univariate linear mixed model per variant per phenotype, to estimate the variant effect sizes, standard errors, and p -values while adjusting for population structure via eigen decomposition of the genomic relatedness matrix. p -values were Benjamini-Hochberg corrected to control the false discovery rate.

Variant effect sizes are usually measured with considerable error in GWAS, and univariate GWAS ignores the correlations that likely exist between effects on the four phenotypes we measured. Furthermore, a key aim of our study is to estimate the frequency of sexual-antagonistic and age-antagonistic variants. The analysis using **mashr** can... and it can be valuable to run additional analyses on the effect sizes to estimate the true effect sizes, and to identify which effects are robust and mitigate ‘winner’s curse’ effects (REF).

We estimated selection on each variant using linear regression (REFS). The statistical model used to test each variant was a simple linear model with the formula $Y \sim \text{Genotype}$, where Y is relative fitness (e.g. the predicted line mean for male early-life fitness, divided by the average of all the predicted line means), and Genotype is the genotype of the focal line. Genotype was coded as a 0 for lines homozygous for the reference allele or 1 for lines homozygous the alternate allele for the focal variant - heterozygous loci are rare in the DGRP, and were excluded from analysis (because the genotype is unknown for these loci in our study). We defined the reference allele as the one that was most common across the entire panel of DGRP lines ($n = 205$), such that a positive slope means that the minor (i.e. rarer) allele confers higher fitness, and a negative slope means that the major (commoner) allele confers higher fitness. For each variant, we recorded the effect size, the associated standard error, and the t , df and p values for the test.

Our approach is equivalent to performing a GWAS with relative fitness as the response variable. The reason we did not use the GWAS pipeline generously provided online by the creators of the DGRP is that we wished to obtain effect size for every variant, and the pipeline only provides effect size for statistically significant variants. However, we did compare the results of our analysis with the results obtained by the Mackay lab’s pipeline, and obtained essentially identical results (e.g. our analysis identified the same statistically significant variant), suggesting that our approach were very similar. One difference is that the Mackay lab’s pipeline estimates the effects of each variant after correcting for the presence/absence of Wolbachia and chromosomal inversions in each line. We conducted pilot analyses which showed that including these variables yielded very similar results (because Wolbachia and chromosomal inversions were both unassociated with fitness in our study; $p > 0.05$), and so we elected to leave them out of our models for simplicity.

Measuring the specificity of selection across sexes and age classes

We calculated a selection index, termed I , for each variant, by adapting the formula from Innocenti and Morrow (XX):

FORMULA

Innocenti and Morrow referred to I as an “index of sex-specific selection”, but I is equally useful for any study that measured selection in two categories of individuals; in our case, we calculated I to compare the effects of each variant on A) male and female fitness, and B) early life and late life fitness.

When I is positive, selection is “concordant”, meaning that there is selection in both sexes or age classes, in the same direction (i.e. the same variant is associated with elevated fitness in both cases). When I is close to zero, selection is absent in one or both cases. When I is negative, selection is “antagonistic” meaning that there is selection in both sexes or age classes, but the variant that is associated with higher fitness is different in each case. We calculated I to compare four kinds of selection: A) males vs females, in early life; B) males vs females, in late life; C) young vs old males; D) young vs old females.

To numerically estimate the uncertainty associated with each estimate of I , we generated 1000 independent samples of b_i and b_j by drawing random numbers from a normal distribution with a mean and standard deviation obtained from the models used to estimate b_i and b_j . This yielded 1000 estimates of I , from which we recorded the median and 95% quantiles, which approximate the 95% confidence limits on I given the uncertainty associated with b_i and b_j .

Annotations for each variant

We relied on the annotations generated by the creators of the DGRP, who used the software SnpEff to classify each variant by site class (e.g. whether the variant is in an intron, or a non-synonymous codon position, etc). Additionally, we assigned a list of KEGG and GO terms to each variant, which matched those associated with the gene (or genes) in which that variant resides (obtained from NCBI).

Results

Variance and covariance in fitness across lines

There was substantial variation in line mean relative fitness, for both sexes and both age classes (Figure 1). All correlations in Figure 1A are positive (Table SX), indicating that lines with high male fitness tended to have high female fitness, and lines with high early-life fitness tended to have high late-life fitness. Despite the overall positive correlations, some lines had higher than average female fitness but below-average male fitness, and some that ranked highly for early life fitness had low late-life fitness (and vice versa; Figure 1B). In sum, the data suggest that the majority of genetic variance in fitness is concordant across sexes and age classes, but provide support for the existence of alleles with antagonistic effects on the four phenotypes (particularly male and female fitness).

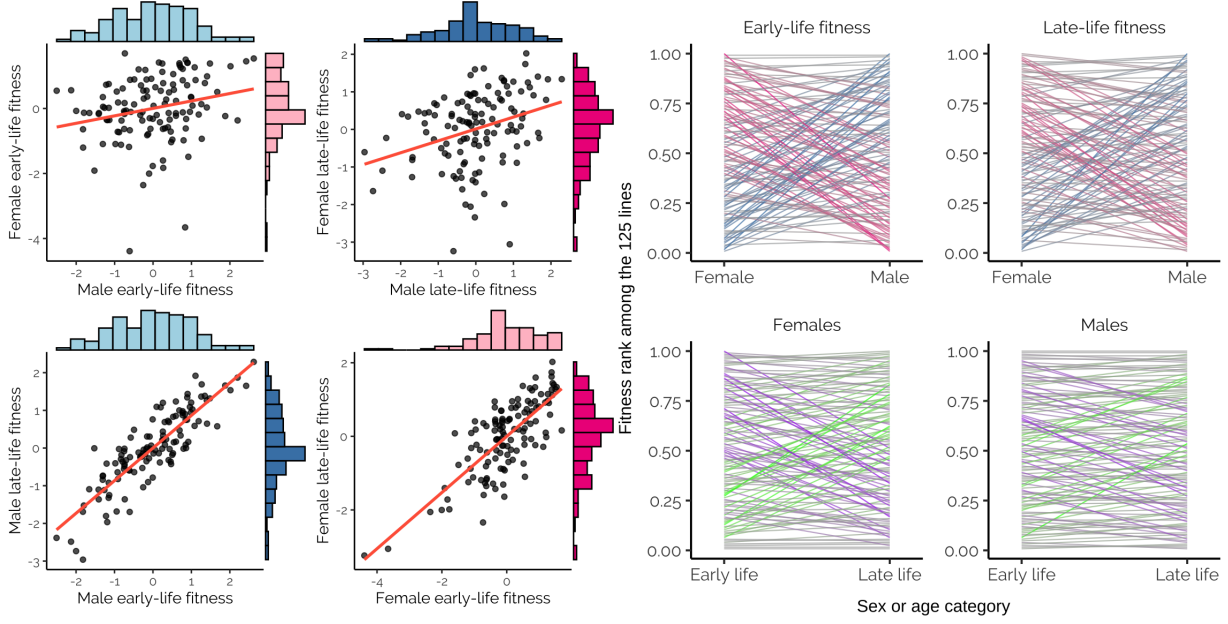


Figure 1: Panel A shows correlations among estimated line means for the four fitness components from Bayesian mixed models adjusting for block effects. Panel B shows the same data expressed as ranks, and the panels illustrate how the ranks co-vary between pairs of phenotypes (the colours indicate the slope of the line, e.g. redder lines have high-ranked female fitness and low-ranked male fitness).

Genetic variance and covariance in fitness

The estimated G matrix for our four traits is shown in Table 1. Fitness was highly heritable...

Distribution of fitness effects across variants

Figure 2 plots the mashr-adjusted effect sizes of the polymorphic loci on each of the four phenotypes (see Table SX for associated Spearman correlation statistics). As expected based on the line means in Figure 1, there was a positive correlation, such that at loci where the minor allele increase or decrease female fitness, the minor allele tended to have a similar effect on male fitness. The correlation was similarly strong in both age classes, though the slope was shallower in the late life age class due to a smaller effect sizes on male fitness as compared to the early life age class.

The mean effect on fitness across all variants was significantly negative for all but one of the phenotypes, which means that the minor alleles were, on average, associated with lower fitness than the major alleles (Table SX). The exception was male early-life fitness, for which the reverse was true: the major allele tended to be associated with lower fitness than the minor allele. However the average variant effect size differed by only a small amount from zero (i.e. the value expected under the null hypothesis that the minor and major alleles are equally likely to be the beneficial one).

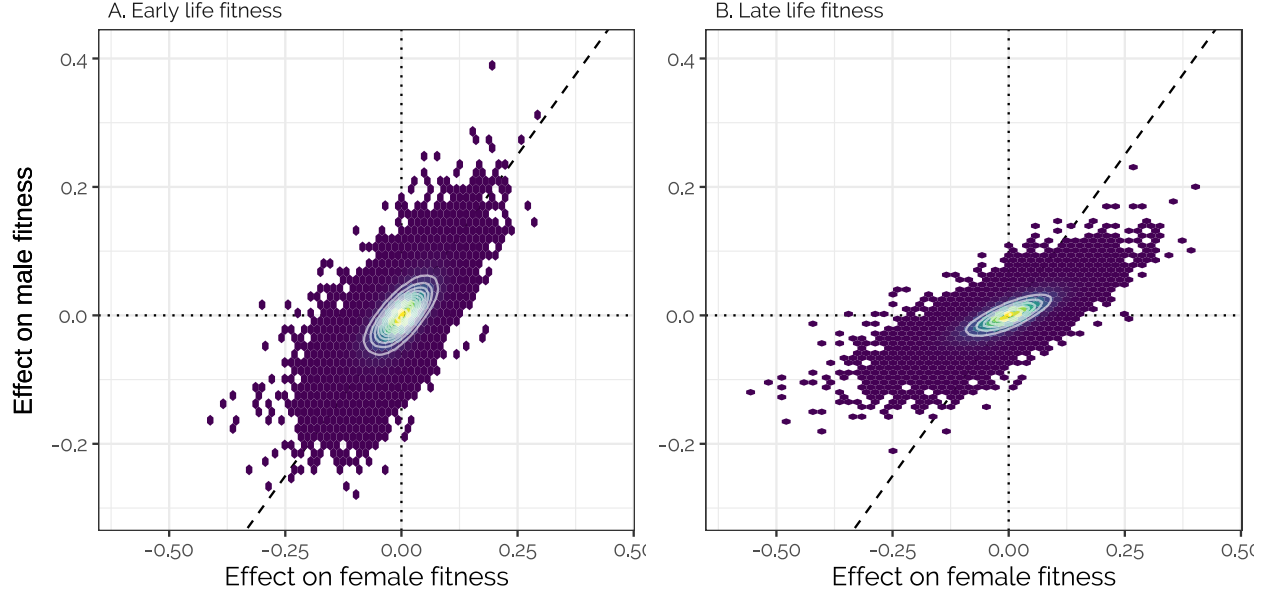


Figure 2: Effect sizes of 1,207,357 loci (i.e. groups of one or more polymorphic sites in complete linkage disequilibrium) on male and female fitness, plotted separately for the early-life and late-life estimates. The effect sizes estimated using GEMMA have been corrected using mashr, using the data-driven method to apply shrinkage (Figure SX shows the raw estimates). The data have been binned into hexagons, with the colour and contour lines indicating the number of loci. The diagonal line represents $y = x$. Positive effect sizes indicate that the minor allele is associated with higher fitness.

Top hits from the GWAS

We used the SNP clumping method in PLINK to identify groups of loci in linkage disequilibrium with one another, for which at least one group member affected one of the four phenotypes with $-\log_{10}(p) > 5$ (Table SX). There were 83 such groups of loci (15 groups with $-\log_{10}(p) > 6$, and 3 with $-\log_{10}(p) > 7$). Many of these regions overlapped genes. For example, we found a significant variant ($-\log_{10}(p) = 7.3$) in an intron of *serotonin receptor 1A* (FBgn0004168), where the minor allele (MAF = 0.11) reduced fitness in both sexes. A complete list of affected genes is given in Table SX.

We found one polymorphic region that had a sexually antagonistic effect on fitness (with the major allele benefitting females and harming males; MAF = 0.10), overlapped six genes: *Death regulator Nedd2-like caspase* (*Dronc*; FBgn0026404), *variable nurse cells* (*vnc*; FBgn0263251), and *defective proboscis extension response 6* (*dpr6*; FBgn0040823) plus 3 uncharacterised genes (FBgn0036062, FBgn0036063, and FBgn0259932). *Dronc* mediates caspase-dependent cell death, and is implicated in the DNA damage response, sperm differentiation, and cell specification, while *vnc* is involved in histone acetylation and oogenesis. *dpr6* is involved in synapse organisation and the perception of chemical stimuli; another *dpr* gene is required for proper male courtship (<https://journals.plos.org/plosgenetics/article?id=10.1371/journal.pgen.0030216>). The three uncharacterised genes all show high expression in the ovary and less expression in the testis (modENCODE Tissue Expression), implying roles in reproduction.

Top hits from the TWAS

Almost all loci affect fitness, or are close to a locus that does

Inspired by Figure 1C in Boyle et al. (REF), we sorted all of the variants by their fitness effects, placed them in bins of 1000, and then calculated the average fitness effect for each bin. Figure 3 shows that there was a very tight correlation (XXX) between the average effects of the variants in each bin on male and female fitness, and that the shape of the relationship was different between age classes. For the early life age class, there was a very tight positive correlation between the effect size on males and females for loci where the minor allele increased fitness in females (i.e. those $x > 0$ in Figure 3A). For loci where the major allele increased female fitness ($x < 0$ in Figure 3A), the correlation was non-linear, such that the effect of the

major allele on male fitness became less positive as its effect on females increased. Figure 3A thus implies a shortage of loci where the major allele is strongly beneficial to fitness in both sexes in the early life age class (though there was no shortage of loci where the minor allele was strongly beneficial in both sexes). Figure 1B showed similar patterns, except that the relationship did not change between loci where the minor versus major allele benefitted females.

Figure 3 reaffirms the results shown in Figures 1-2 that there is a positive genetic correlation between male and female fitness, and that this correlation changes between age classes, but Figure 3 also implies that our four phenotypes are highly polygenic (or ‘omnigenic’; Boyle et al. REF) and are affected by large numbers of loci with small effects. The male and female fitness measurements were collected independently, and so Figure 3 allows us to distinguish small but genuine effects from statistical noise, despite the low power of our study (and most GWAS) to detect variants with weak effects. To see why, consider an alternative hypothesis, in which the great majority of variants have no effect on fitness, and the genetic (co)variance in phenotypes observed across lines (Figure 1 and QUANT GEN TABLE) resulted from comparatively few (e.g. dozens) of variants with comparatively large effects. The plot in Figure 3 would then be flat in the centre with steep inflections at each end. The straight line that we see instead suggests that there is a large number of variants that each affect fitness - typically in both sexes, in the same direction - whose effect sizes range from tiny to moderate. Put another way, the effect size of each 1000-variant bin on female fitness was replicated for male fitness, suggesting that many of the non-zero effect size estimates reflect real non-zero effects, as opposed to statistical sampling variance.

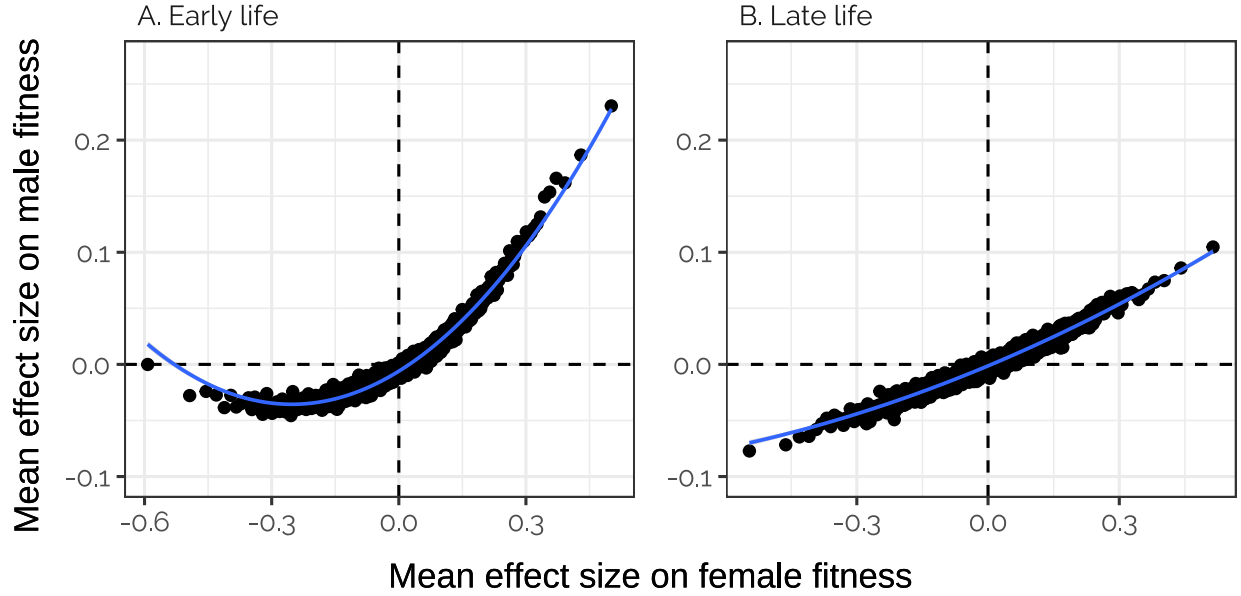


Figure 3: Mean effect size on each phenotype for groups of 1,000 loci. The loci were first sorted in order of their effect on female fitness, then binned and their effect sizes on each sex were calculated. The line shows a quadratic linear regression fit with standard error in grey.

Mixture modelling results from the GWAS and TWAS

Figure 4A shows the results of a mixture modelling analysis of the GWAS results, from the R package `mashr` (in canonical mode). `mashr` allows one to estimate the proportions of various types of locus, using a set of covariance matrices (4×4 matrices, in our case) that are specified *a priori* by the user. To be inclusive of all possibilities, we hypothesised that loci might A) have no effect on any of the four phenotypes, B) affect one of the four phenotypes and not the others, C) affect both phenotypes for one sex, either concordantly or antagonistically across age classes, D) affect both phenotypes for one age class, either concordantly or antagonistically across sexes, or E) affect all four phenotypes in the same direction or differing directions (including complex types, such as loci that are sexually antagonistic in one age class and sexually concordant in the other). Several of these possibilities were calculated to apply to fewer than 1% of loci by `mashr`, and thus are not plotted in Figure 4. We ran `mashr` once for all the loci, and ran it a further 5 times using just the loci from one of the 5 major chromosome arms.

Figure 4A shows that the most common type of locus was estimated to be the type that has an identical effect size on all four phenotypes (recorded as ‘sexually concordant’ in Figure 4A). The model estimated that 0% of loci have age-specific or age-antagonistic effects, and so these types are not plotted. Loci with no effect on fitness were rare, despite the model’s prior that this type of locus is ten-fold more common than any other type, again indicating that most loci either affect fitness or are closely linked to a locus that does. Loci with sex-specific effects were estimated to be quite common, especially loci with sexually antagonistic (12%, genome-wide) or female-specific effects (19%). The X chromosome was inferred to have the highest proportion of sexually antagonistic loci (24%), with an estimate that was twice as high as the estimate for the whole genome, and substantially higher than for other chromosomes (e.g. 7.4% for chromosome 2L).

Distribution of fitness effects across chromosomes

To do...

Discussion

Tables

Table 1: List of variables, and their corresponding parameter(s) in the model, which were varied in order to study their effects on extinction. `# {r xtable, results="asis"} # print_table1() #`

Figures